

KIM JULONGBAYAN

AI Systems and Automation Engineer | Agents, Integrations and Business Process

Quezon City, Metro Manila, Philippines | kljulongbayan@gmail.com | +63 961 529 9305

Portfolio: kimlj.dev | GitHub: github.com/kimlj | LinkedIn: linkedin.com/in/kimjulongbayan

PROFESSIONAL SUMMARY

- I build the systems that take manual work off a team: AI agents on the Claude and OpenAI APIs, automations across n8n, Make, Zapier and Power Automate, and the APIs and databases underneath them. I am also the person who goes looking for the work worth automating rather than waiting to be handed it.
- 776 hours and 16,365 prompts in Claude Code across 43 projects (150 active days, as of Sep 2026), counted from session transcripts and published at kimlj.dev.

TECHNICAL SKILLS

- AI agents and LLM APIs: Claude and OpenAI, tool use, prompt engineering, checkable structured output, agents that check each other's work, retrieval over a knowledge base, MCP in daily use, per-call cost tracking
- Automation platforms: n8n (self-hosted), Make, Zapier, Microsoft Power Automate, Google Apps Script, Looker Studio, scheduled jobs, conditional routing, human approve and reject steps
- Integrations: REST APIs, webhooks, JSON, OAuth, Google Workspace, Slack alerts, Telegram Bot API, Supabase, third-party APIs behind rate limits
- Data and languages: Python, JavaScript, TypeScript, SQL; PostgreSQL with row-level security, Supabase, MySQL, SQLite, schema design and versioned migrations, embeddings and similarity search, automated KPI reporting
- Keeping it running: error handling and retries, alerting on failure, Sentry, Docker on a Linux VPS, Git, and documentation so a team can maintain it without me

EXPERIENCE

MDS Pro Solutions - Full Stack Developer (Contract) | Jun 2026 - Present

- Sole engineer on the timekeeping and billing platform a US company runs on: 15 people use it every workday, real payroll and invoices come out of it, nearly 400 commits since June.
- Built an AI agent into it on the Claude API: staff ask a question in plain English, the model turns it into a database query, and a locked-down account runs it that can only ever read. Every answer records what it cost.
- Connected the platform to the owners' Google Sheet over the Sheets API: a scheduled sync fills blank cells only and flags anything that disagrees, so an automation can never overwrite a figure a person typed.
- Designed the approval path deliberately: records are append-only, a correction is requested, a human approves, and the result is stamped with who changed it. Automation proposes, a person decides.

Independent AI and Automation Developer | 2018 - Present

- Built and shipped automation and AI systems end to end, including the integrations, the monitoring and the handover documentation.

SELECTED SYSTEMS

Shift-Ops Console - One Rules Engine Across Six Platforms

- An engine that flags attendance and compliance exceptions across a simulated 100-person workforce: n8n runs the rules, Make sends the daily summary, Zapier fires urgent Slack alerts, Power Automate runs the approve and reject loop, Looker Studio draws the KPI dashboards.
- Then I measured it rather than trusting it: 370 planted problems, 359 caught, 97 percent. The measurement also caught my own scoring being wrong, which was the more useful result.

Jobsift - An Agentic Pipeline That Grades Its Own Output

- Four roles instead of one call: an extractor that never sees the candidate, a drafter, a verifier run once per claim against the trusted sources, and a supervisor that sends unsupported claims back to be redone.
- Separately, a knowledge agent live on kimlj.dev answers in my voice from a knowledge base generated out of the site itself, so the page and the agent can never say different things.

EDUCATION

- Bachelor of Science in Information Technology, AMA Computer College Makati (3 of 4 years completed)